

Edoardo Debenedetti

PHD STUDENT IN CS @ ETH ZÜRICH

☎ (+41) 76 699 43 27 | ✉ edebenedetti@inf.ethz.ch | 🌐 <https://edoardo.science> | 💻 dedeswim | 📞 0000-0003-3343-9477 | 📧 Edoardo Debenedetti

Education

ETH Zürich - Federal Institute of Technology Zürich

Zürich, Switzerland

PHD IN COMPUTER SCIENCE

08/2022 - 12/2026 (exp.)

- Focus: **Real-world machine learning security and privacy**, advised by **Prof. Florian Tramèr** in the **SPY Lab**.
- **IT Coordinator** for the group: managing the GPU servers and hardware resources.
- Fully funded by the **CYD Doctoral Fellowship**, awarded by the Armasuisse Cyber-Defense Campus.

EPFL - Federal Institute of Technology Lausanne

Lausanne, Switzerland

MSC IN COMPUTER SCIENCE

09/2019 - 04/2022

- **GPA 5.63/6**, focus on **Machine Learning** ∩ **Security** ∩ **Privacy**.
- Master's Thesis about the **adversarial robustness of Vision Transformers** supervised by **Princeton University's Prof. Mittal**.

Politecnico di Torino

Turin, Italy

BSC IN COMPUTER ENGINEERING

09/2016 - 07/2019

- **GPA 28.4/30**, graduation mark 110/110, **top 9%**.
- **Exchange year at 同济大学** (Tongji University), in Shanghai (China), supported by a **full scholarship** granted to the top 31% applicants.

Experience

Bloomberg LP

London, United Kingdom

SOFTWARE ENGINEERING INTERN

07/2021 - 09/2021

- Worked in the **Multi Asset Risk System** team, on the re-design and implementation of the configuration of a distributed logging library.
- Move the configuration of a **distributed logging library** from an internal technology to a **centralized SQL DB**, using a **cache** and a **C++ service**.
- The configuration is checked **~1M times per minute**, and the usage of the cache gave a **~23x speed improvement** w.r.t. querying the DB.

armasuisse Cyber-Defence Campus

Lausanne, Switzerland

RESEARCH INTERN

08/2020 - 01/2021

- Worked on **Machine Unlearning** and **Membership Inference Attacks** against Generative Models, supervised by **Prof. Mathias Humbert**.
- Adapt the **MIA** technique proposed by the *GAN-Leaks* work (by Chen et al.), to work after the removal some datapoints from the training set.
- The technique achieved **promising results** when attacking DCGAN trained on the CelebA dataset

Conference papers

- **Debenedetti, E.**, Sehwal, V., Mittal, P., "*A Light Recipe to Train Robust Vision Transformers*", First IEEE Conference on Secure and Trustworthy Machine Learning, February 2023.
- Croce*, F., Andriushchenko*, M., Sehwal*, V., **Debenedetti*, E.**, Flammarion, N., Chiang, M., Mittal, P., Hein, M., "*RobustBench: a standardized adversarial robustness benchmark*", Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track, 2021. (* equal contribution). A preliminary version appeared at the ICLR 2021 Workshop on Security and Safety in ML Systems.

Manuscripts and workshop papers

- **Debenedetti, E.**, Severi, G., Carlini, N., Choquette-Choo, C. A., Jagielski, M., Nasr, M., Wallace, E., Tramèr, F., "*Privacy Side Channels in Machine Learning Systems*", arXiv ePrint 2309.05610.
- **Debenedetti, E.**, Carlini, N., Tramèr, F., "*Evading Black-box Classifiers Without Breaking Eggs*", 2nd ICML Workshop on New Frontiers in Adversarial Machine Learning, 2023. **Oral presentation**.

Honors and Awards

- 2023 **Oral presentation - AdvML Frontiers Workshop**, Top 10% accepted papers.
- 2023 **CYD Doctoral Fellowship**, full PhD funding for 4 years, worth **USD 516'000** (CHF 461'000), from Armasuisse CYD Campus and EPFL.
- 2021 **Google TPU Research Cloud Program**, extensive **hardware support for 8 months** to work on the Master's Thesis.
- 2021 **Best Paper Honorable Mention - ICLR Workshop on Security and Safety in ML Systems**, top 2 out of 50 accepted papers.

Teaching

Classes

- **Information Security Lab** – ETH Zürich: 2022, 2023 (Teaching Assistant)
- **Large Language Models** – ETH Zürich: 2023 (Teaching Assistant)

Service

Reviewer

- **NeurIPS Datasets and Benchmarks Track**: 2022, 2023
- **CCS AI Sec workshop**: 2023

Open Source Maintainer

- **RobustBench**: adversarial robustness benchmarking library and model zoo.
 - More than 150 models spanning 3 datasets and 3 threat models.
 - **517 stars**, with 379 unique cloners in 2 weeks (measured in September 2023).
 - Refactored the code to improve the extensibility of the library.

Repository at <https://github.com/RobustBench/robustbench>.

Conference service

- **Competition organizer at SaTML 2024**: co-organizing the *Large Language Models Capture-the-Flag*. More than **70 teams and 100 users** signed up.
- **Volunteer at NeurIPS 2021**: helped with monitoring the website and technical issues.

Invited talks

- **ACL SIGSEC** – *Privacy Side-channels in Machine Learning Systems*, 2023.
- **TU Graz EfficientML Reading Group** – *Privacy Side-channels in Machine Learning Systems*, 2023.